

pythonCreditEDA

March 22, 2018

0.1 Credit EDA

```
In [3]: #Packages
        #for data manipulation and operations
        import pandas as pd

        #for plotting
        from matplotlib import pyplot as plt
        %matplotlib inline
        import seaborn as sns

        #for statistical modeling
        import statsmodels.api as sm
        import numpy as np
```

```
C:\Users\Matt\Anaconda3\lib\site-packages\statsmodels\compat\pandas.py:56: FutureWarning: The pandas.core
from pandas.core import datetools
```

```
In [4]: #read in the data and verify
        data = pd.read_csv('creditdata.csv')
        data.head()
```

```
Out[4]:
```

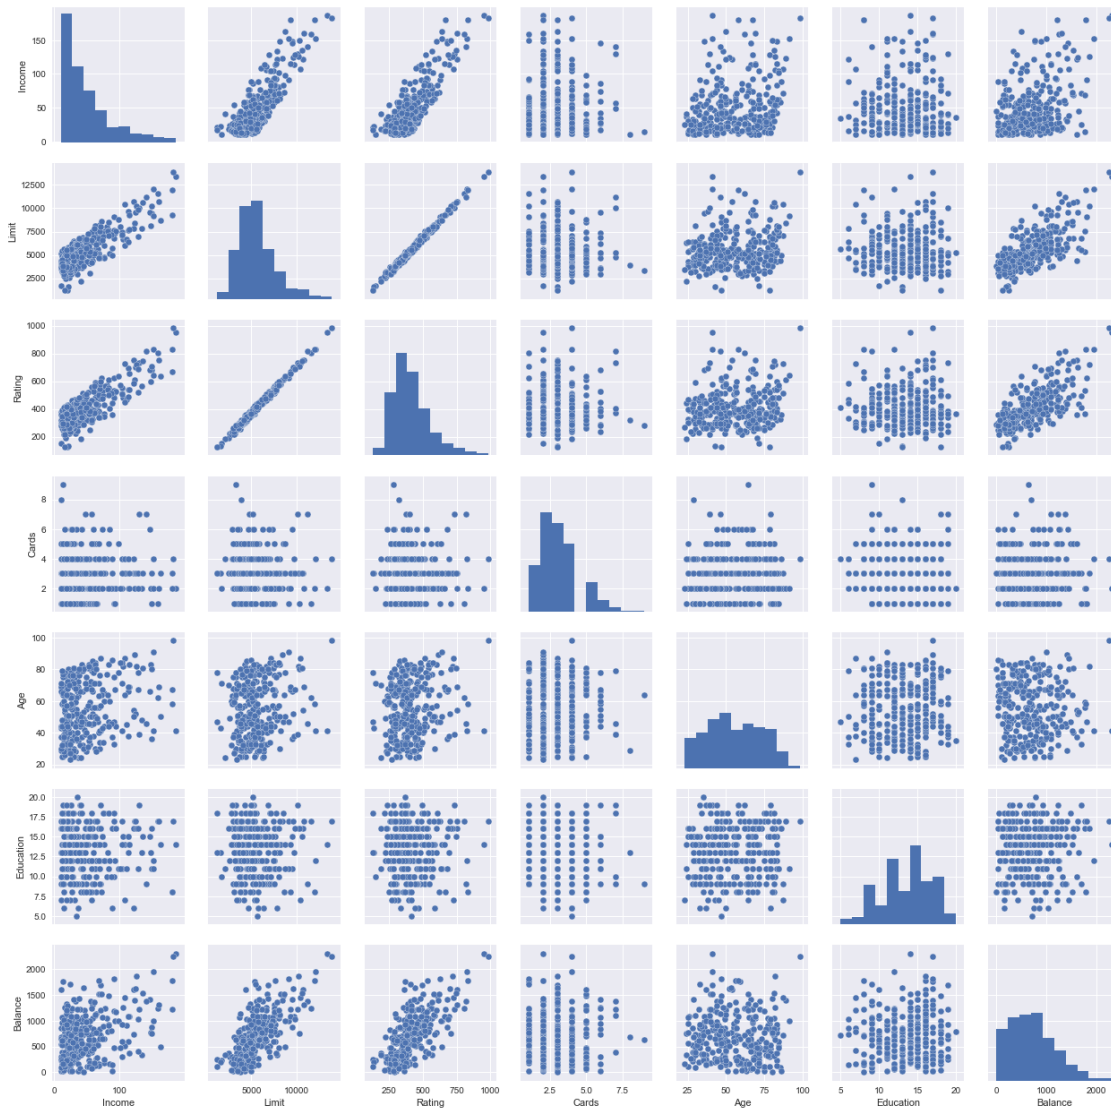
	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	\
0	14.891	3606	283	2	34	11	Male	No	Yes	
1	106.025	6645	483	3	82	15	Female	Yes	Yes	
2	104.593	7075	514	4	71	11	Male	No	No	
3	148.924	9504	681	3	36	11	Female	No	No	
4	55.882	4897	357	2	68	16	Male	No	Yes	

	Ethnicity	Balance
0	Caucasian	581.74
1	Asian	1259.22
2	Asian	301.46
3	Asian	878.14
4	Caucasian	509.92


```
In [5]: # #make histograms of all the variables
        # data.hist(bins=50,figsize=(20,15))
        # plt.show()

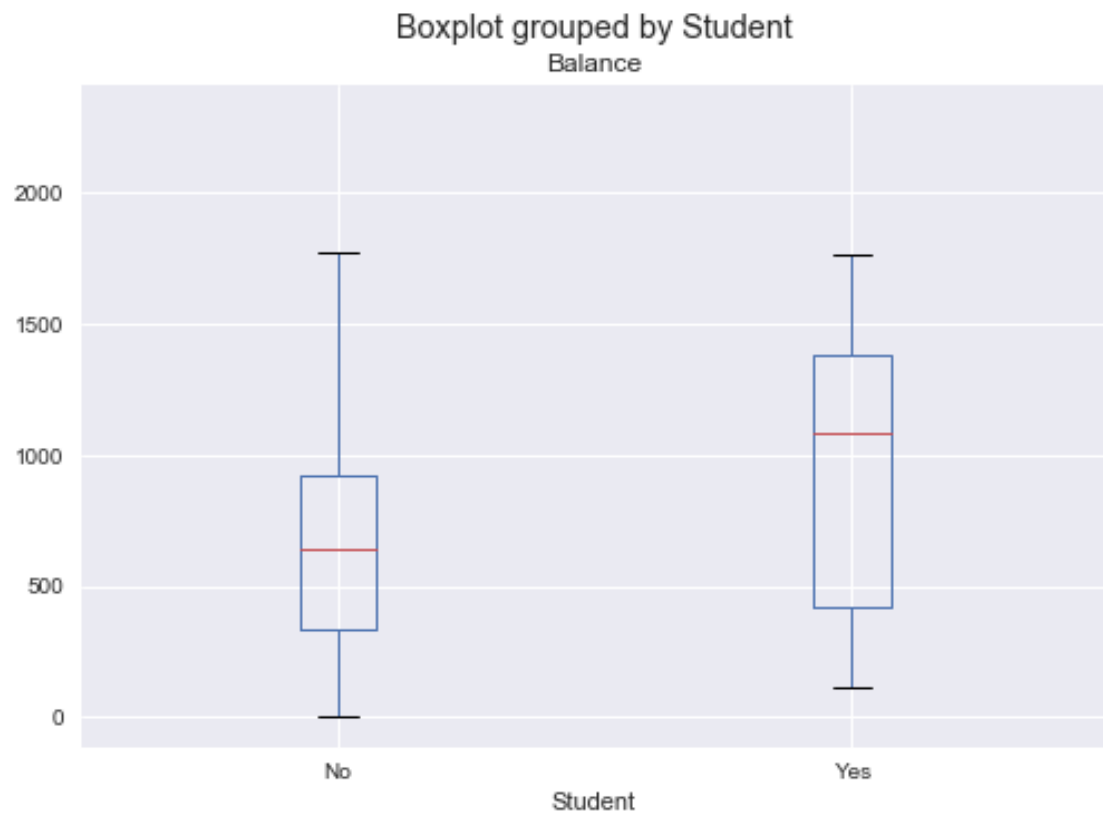
        # g = sns.pairplot(iris, hue="species")
        # sns.pairplot(data, hue='Ethnicity')
        # sns.pairplot(data, hue='Student')
        sns.pairplot(data)
```

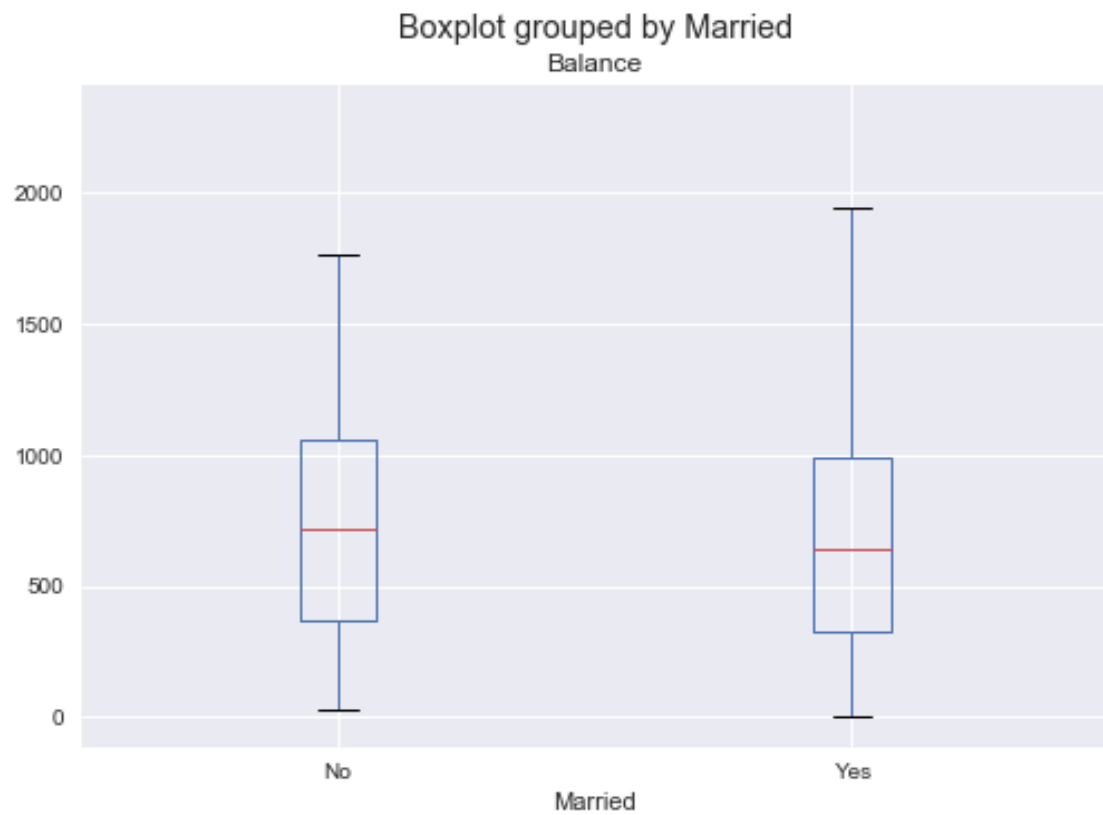
```
sns.despine()
```

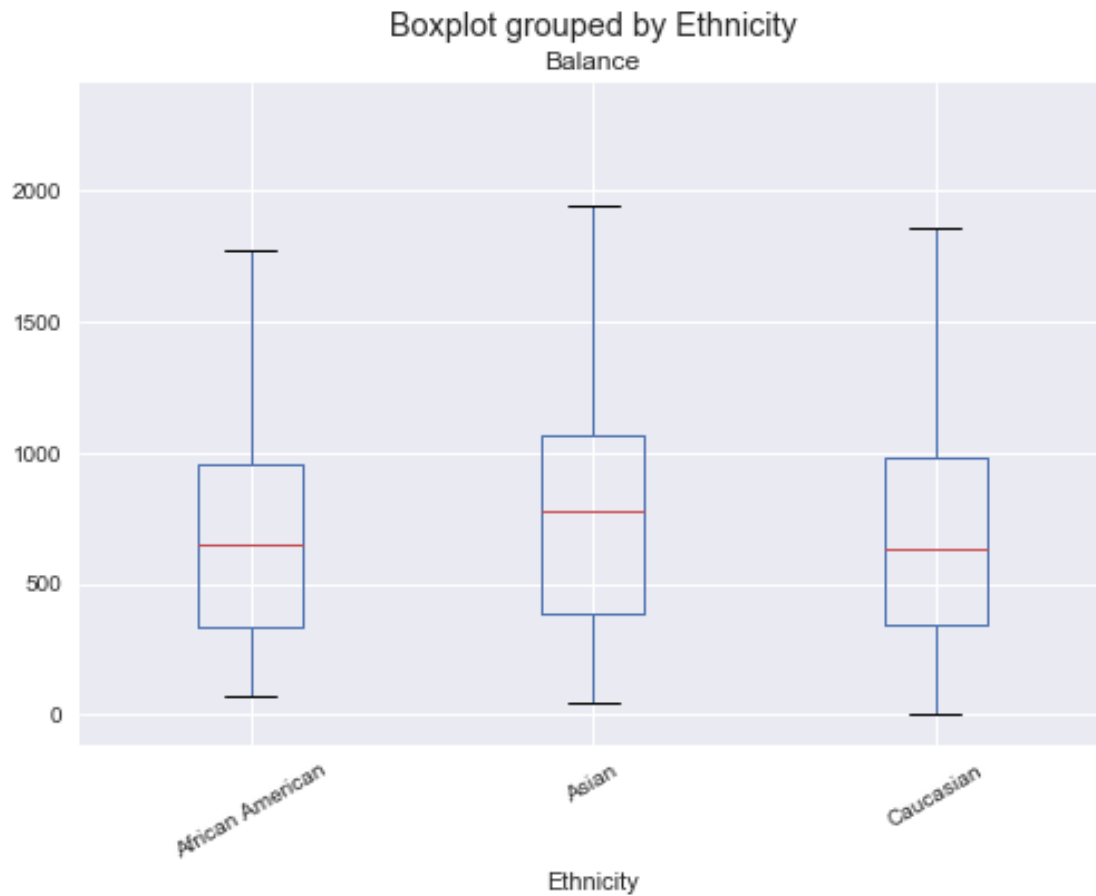


```
In [30]: data.boxplot('Balance',by='Gender')
plt.show()
data.boxplot('Balance',by='Student')
plt.show()
data.boxplot('Balance',by='Married')
plt.show()
data.boxplot('Balance',by='Ethnicity',rot=30)
plt.show()
```






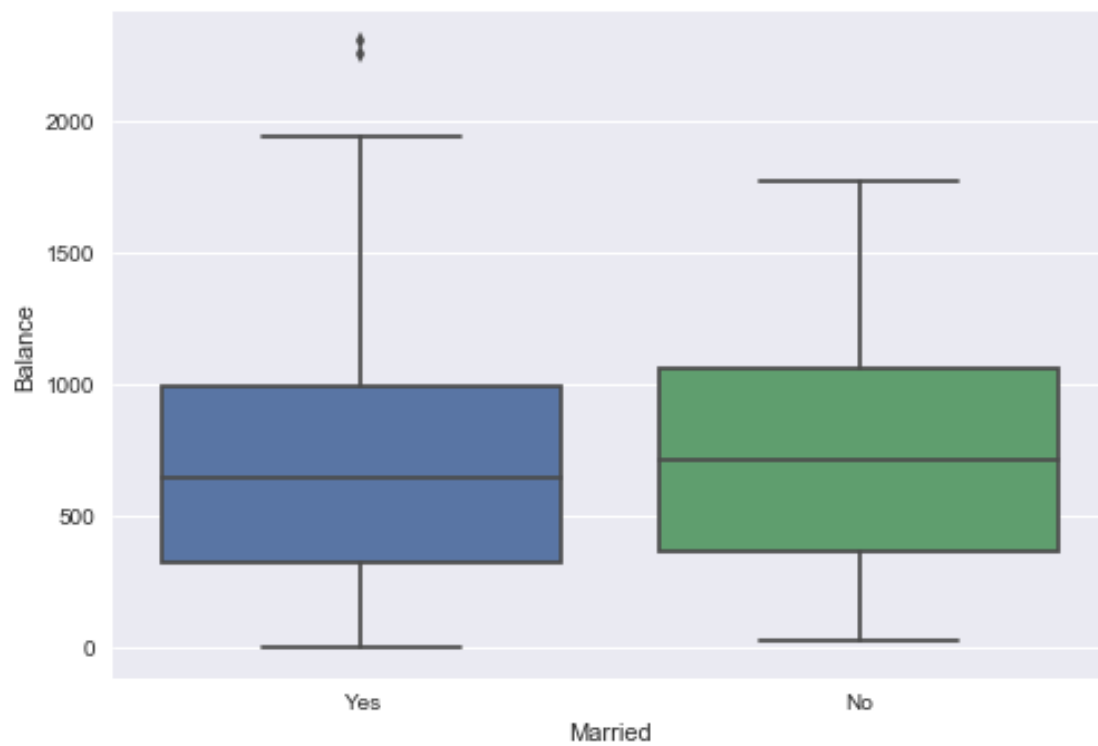






```
In [23]: # ax = sns.boxplot(x="day", y="total_bill", hue="smoker", data=tips, palette="Set3")
sns.boxplot(y='Balance', x='Gender', data=data)
plt.show()
sns.boxplot(y='Balance', x='Student', data=data)
plt.show()
sns.boxplot(y='Balance', x='Married', data=data)
plt.show()
sns.boxplot(y='Balance', x='Ethnicity', data=data)
plt.show()
sns.despine()
```





<matplotlib.figure.Figure at 0x2149d81ccf8>

```
In [7]: data.shape
```

```
Out[7]: (294, 11)
```

```
In [8]: type(data)
```

```
Out[8]: pandas.core.frame.DataFrame
```

```
In [9]: type(data.drop('Balance',axis=1))
```

```
Out[9]: pandas.core.frame.DataFrame
```

```
In [10]: import statsmodels.formula.api as smf
```

```
In [11]: ### basic linear regression (without variable selection)
```

```
# adding '1' as a predictor forces the inclusion of an intercept
# model = smf.ols(formula='DLHRWAGE ~ 1 + DEDUC1 + AGE + AGESQ + WHITEH + MALEH + EDUCH + \
#                               WHITEL + MALEL + EDUCL + DEDUC2 + DTEN + DMARRIED + DUNCOV', data=twinT
model = smf.ols(formula='Balance ~ 1 + Income + Limit + Rating + Cards + Age + Education + \
                    Gender + Student + Married + Ethnicity',data=data).fit()
model.summary()
```

```
Out[11]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

OLS Regression Results

```
=====
Dep. Variable:          Balance    R-squared:                0.787
Model:                  OLS      Adj. R-squared:            0.779
Method:                 Least Squares    F-statistic:          94.99
Date:                  Tue, 09 Jan 2018    Prob (F-statistic):    4.95e-88
Time:                  11:42:02    Log-Likelihood:       -1985.2
No. Observations:      294          AIC:                  3994.
Df Residuals:          282          BIC:                  4039.
Df Model:               11
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-543.4223	90.126	-6.030	0.000	-720.828	-366.016
Gender[T.Female]	-19.7824	24.938	-0.793	0.428	-68.870	29.306
Student[T.Yes]	535.7803	39.204	13.666	0.000	458.610	612.950
Married[T.Yes]	-14.2811	26.093	-0.547	0.585	-65.642	37.080
Ethnicity[T.Asian]	53.3056	36.427	1.463	0.144	-18.398	125.009
Ethnicity[T.Caucasian]	41.4142	30.790	1.345	0.180	-19.192	102.021
Income	-8.4707	0.605	-13.992	0.000	-9.662	-7.279
Limit	0.3681	0.082	4.502	0.000	0.207	0.529
Rating	-0.9379	1.207	-0.777	0.438	-3.313	1.437
Cards	20.5028	10.235	2.003	0.046	0.356	40.650
Age	-2.2444	0.734	-3.056	0.002	-3.690	-0.799
Education	-0.1777	3.957	-0.045	0.964	-7.966	7.611

```
=====
Omnibus:                2.257    Durbin-Watson:                2.041
```

Prob(Omnibus):	0.323	Jarque-Bera (JB):	1.956
Skew:	0.178	Prob(JB):	0.376
Kurtosis:	3.182	Cond. No.	4.39e+04

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 4.39e+04. This might indicate that there are strong multicollinearity or other numerical problems.
 ""

```
In [ ]: ## This one doesn't work
        ## Without a constant
        # import statsmodels.api as sm
```

```
## X = df["RM"]
## y = target["MEDV"]
```

```
X = data[['']]
y = data.Balance
# Note the difference in argument order
model = sm.OLS(y, X).fit()
predictions = model.predict(X) # make the predictions by the model
```

```
## Print out the statistics
model.summary()
```

```
In [1]: pwd
```

```
Out[1]: 'C:\\Users\\Matt\\Documents\\Matt BYU\\10 Winter 2018\\Stat 536\\EDA 1 Credit'
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```